

Container and Kubernetes Networking 101

Arun Sriraman

Platform9 Systems Inc

 @arun_sriraman  sarun87

Before we begin

Slides: <https://goo.gl/6XuEnQ>

Poll



Agenda

Part I - Container Networking 101

- Why container networking?
- Linux namespaces
- container-to-container communication
- Docker n/w'ing - CNM model
- libnetwork & driver types


Part II - Kubernetes Networking 101

- K8s networking fundamentals
- Kubernetes communication
 - Container-to-Container
 - Pod-to-Pod
 - Pod-to-Service
 - Service-to-external
- Container Network Interface
- CNI backend (Flannel, Calico)


Part I - Container Networking

Evolution of infrastructure & networking


Infrastructure



Physical servers




Virtualization (VMs)




Containers

Network



Physical switches and routers



Virtual networks (SDN)



Does virtual networks (SDN) continue to work for containers?


Container networking

- container \longleftrightarrow container communication
- container \longleftrightarrow host communication
- user \longleftrightarrow container communication

Why should we learn the nuts & bolts of the system?

- troubleshooting
- debugging
- building solutions (Service Discovery, Load Balancing)

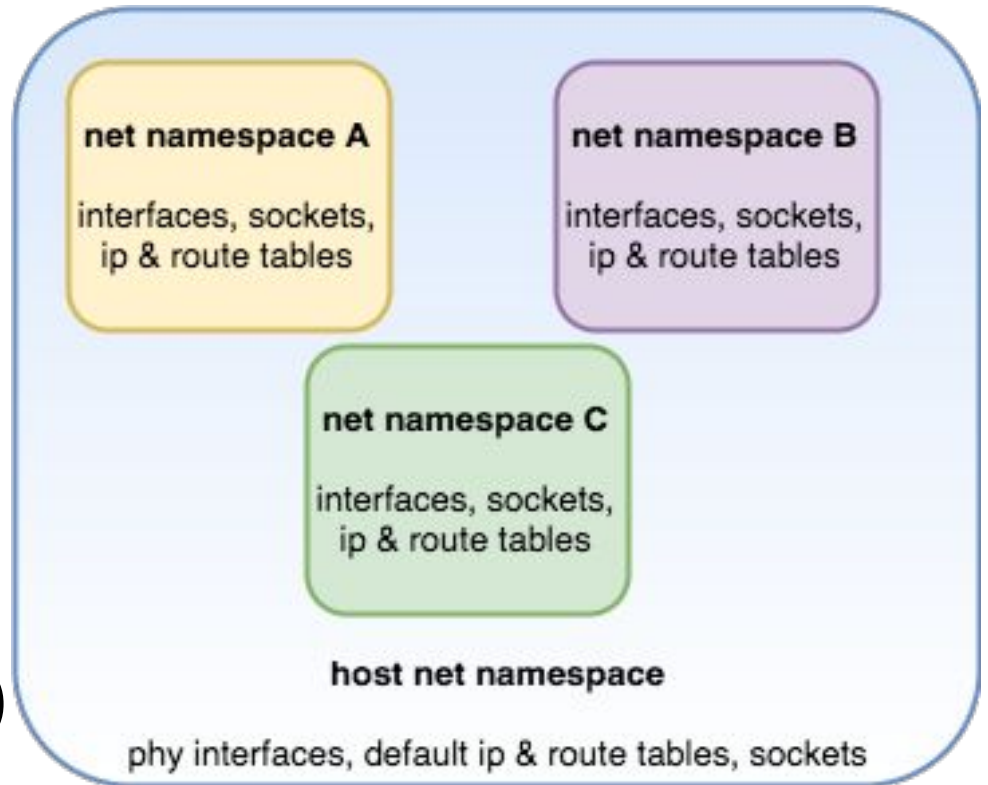
3 steps to container networking



Network Isolation - Network namespaces

A network namespace gets its own private network stack with

- network interfaces (including lo)
- routing tables
- iptables rules
- sockets (ss, netstat)




```
[root@ip-10-0-1-25 ~]# ip netns add A
[root@ip-10-0-1-25 ~]# ip netns add B
[root@ip-10-0-1-25 ~]# ip netns add C
[root@ip-10-0-1-25 ~]# ip netns list
```

Add a network namespace

```
C
B
A
```

List of network namespaces

Namespace A

```
[root@ip-10-0-1-25 ~]# ip netns exec A ip a
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

Host network namespace

```
[root@ip-10-0-1-25 ~]#
```

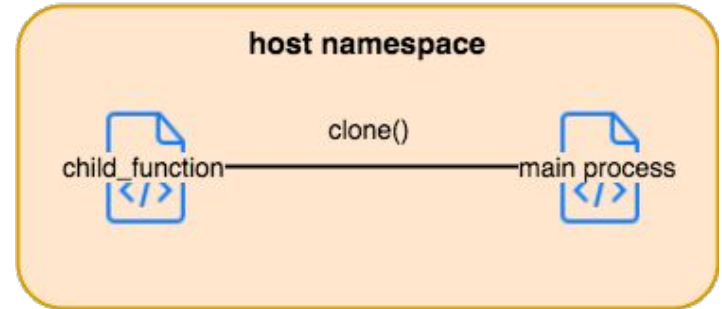
```
[root@ip-10-0-1-25 ~]# ip a
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc pfifo_fast state UP qlen 1000
    link/ether 02:1e:b1:71:b8:6e brd ff:ff:ff:ff:ff:ff
    inet 10.0.1.25/24 brd 10.0.1.255 scope global dynamic eth0
        valid_lft 2573sec preferred_lft 2573sec
    inet6 fe80::1e:b1ff:fe71:b86e/64 scope link
        valid_lft forever preferred_lft forever
```


Container and isolation

```
flags = CLONE_NEWPID |  
        CLONE_NEWNS | CLONE_NEWNET;  
cpid = clone(child_function,  
             childstack,  
             flags, (void *)argv);
```

without `CLONE_NEWNET` flag

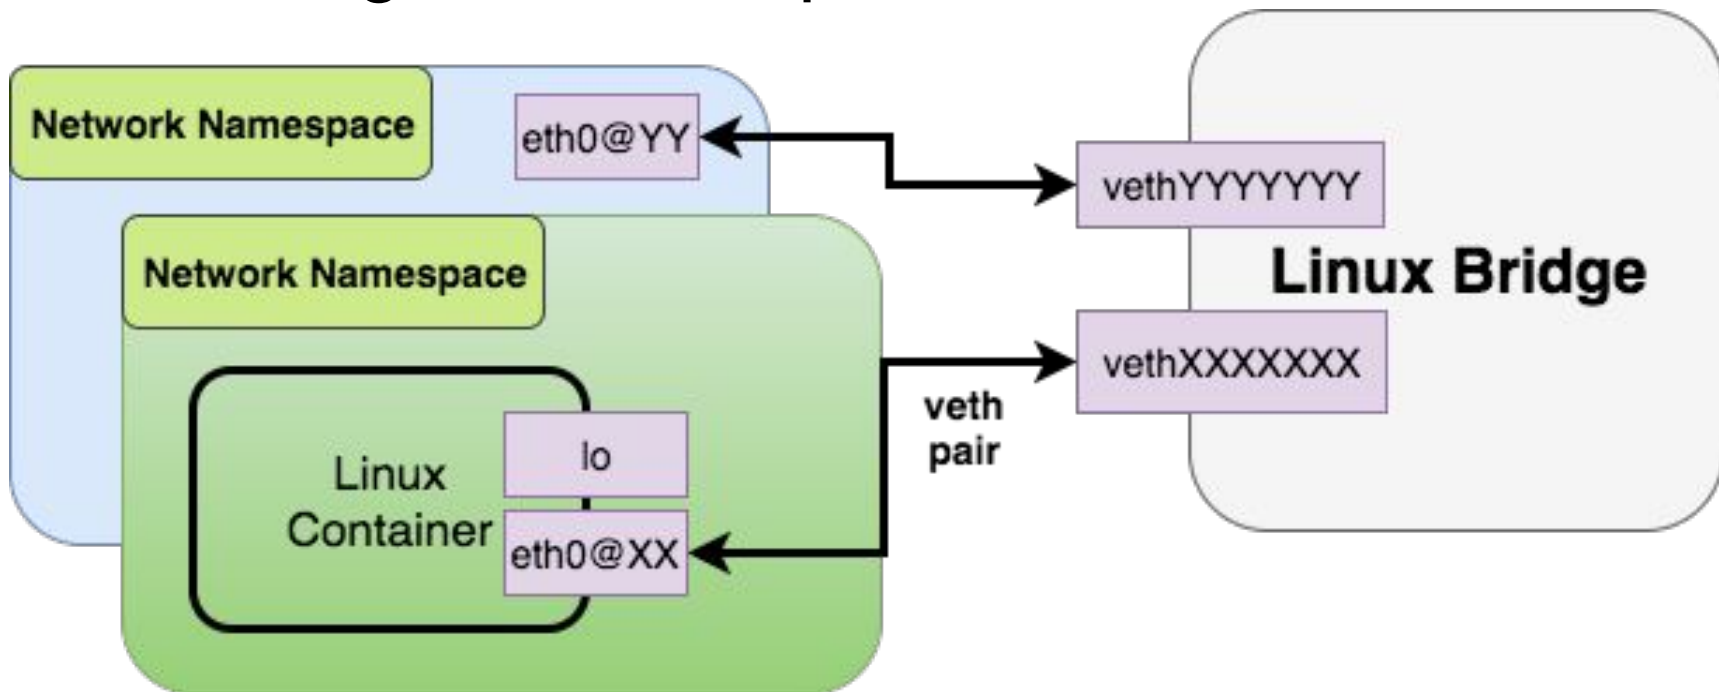


with `CLONE_NEWNET` flag




Communication mechanism

Linux bridge and veth pairs

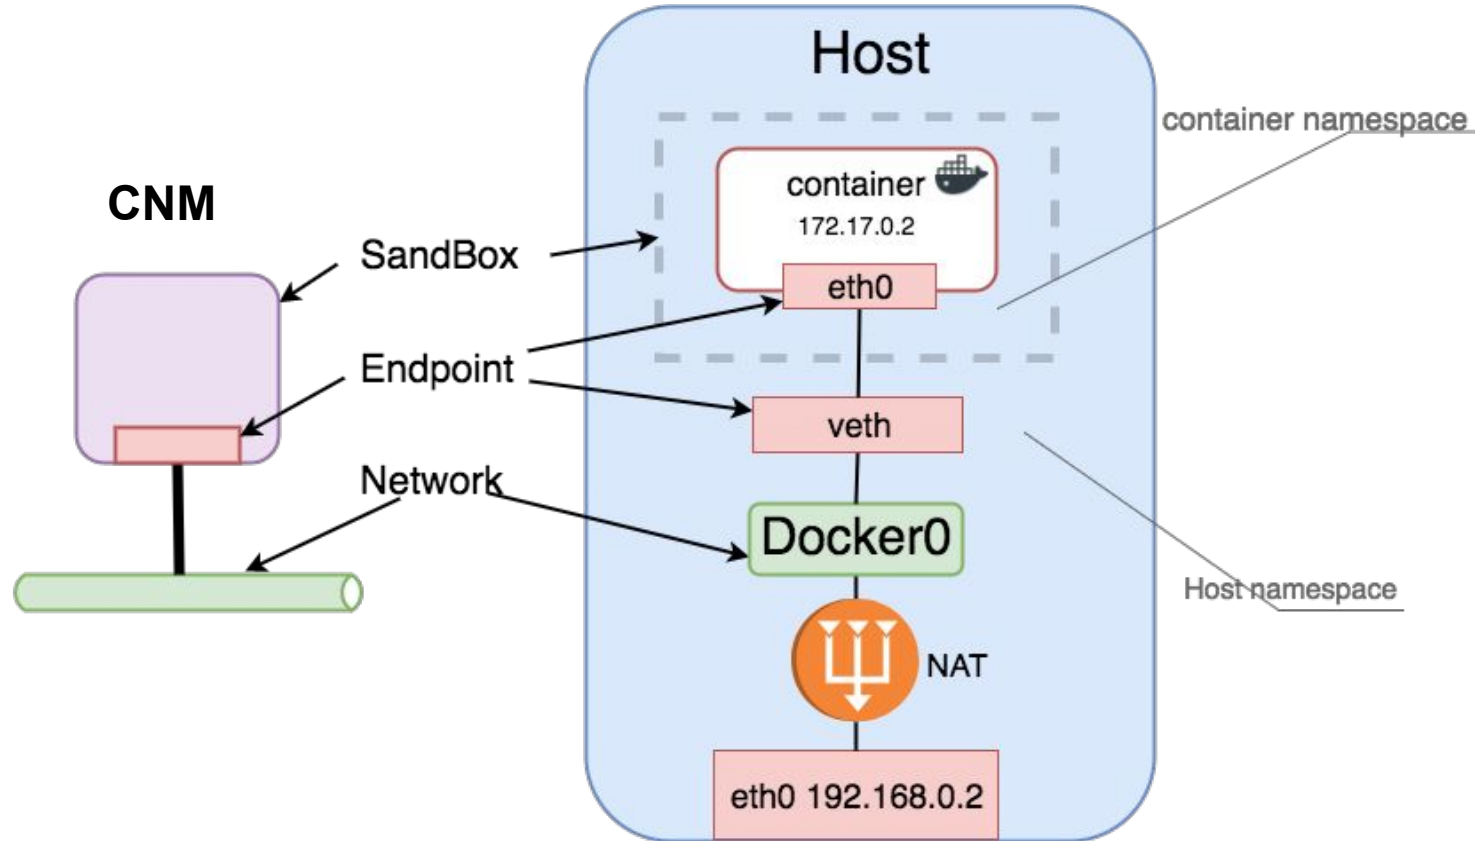


Container Network Model (CNM)

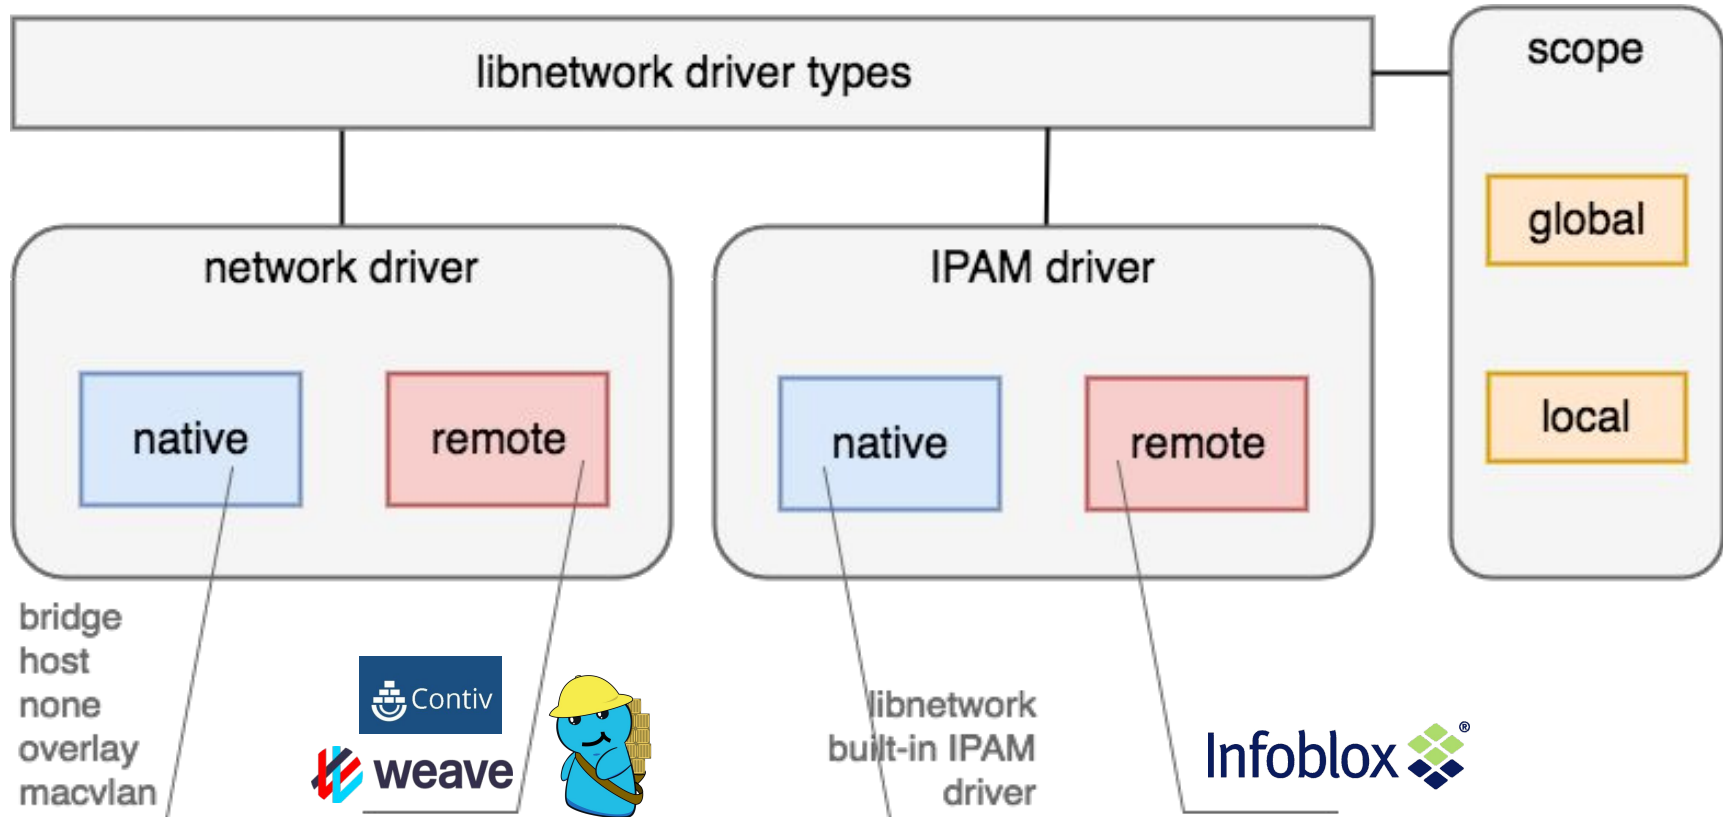
- Project started by Docker
- Separate networking from container runtime as a library
- Components
 - Sandbox
 - Endpoint
 - Network
- Implemented using libnetwork



Mapping CNM to Libnetwork (Docker)



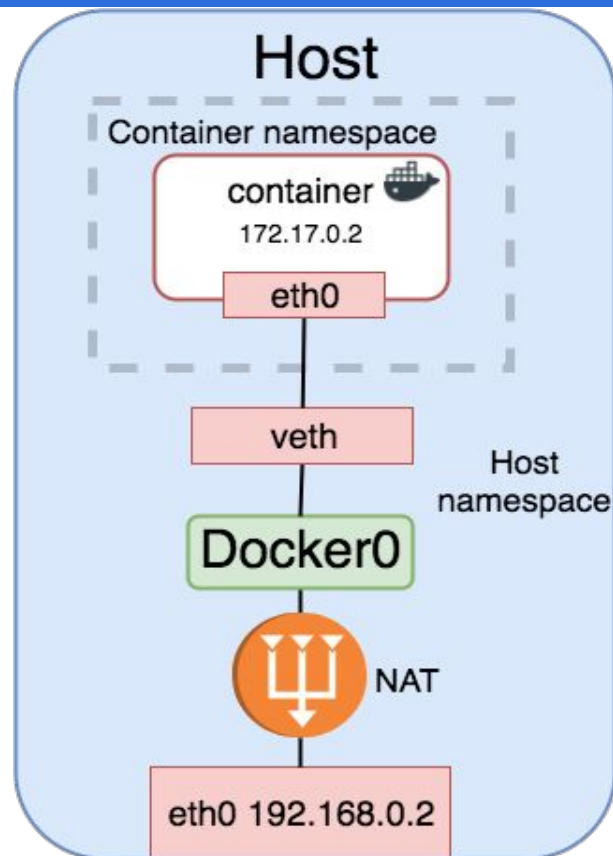
libnetwork driver types



Bridge Driver

- Connects docker containers to the network using a veth pair
- Provides out-of-the-box support for bridge based container networking
- Allows creation of user-defined bridges

```
docker network create --driver bridge  
<name>
```

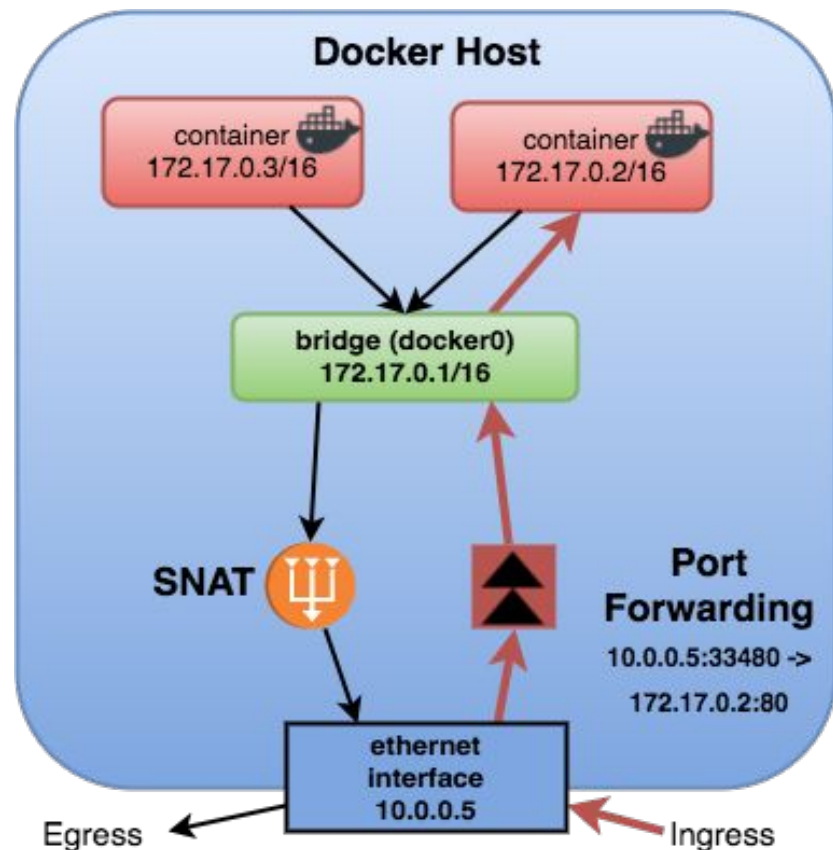


External Access for Containers

- Egress traffic is SNAT'ed using "MASQUERADE" iptable rules
- Ingress traffic is port forwarded in case of bridge driver


```
iptables --list -t nat
```

```
Chain POSTROUTING (policy ACCEPT)
target     prot opt source                destination
RETURN     all  --  192.168.122.0/24      anywhere
base-address.mcast.net/24
RETURN     all  --  192.168.122.0/24      anywhere
255.255.255.255
MASQUERADE all  --  172.17.0.0/16         anywhere
```





Overlay Driver

- Multi-host networking
- First-class citizen in docker networking
- Uses swarm-distributed control plane for centralized mgmt, stability & security
- Uses VXLAN encap (decouples container n/w from physical n/w)
- Overlay datapath entirely in kernel space



Macvlan Driver



Default Networks Created by Docker

'bridge' using **bridge** driver, 'none' using **null** driver, 'host' using **host** driver

```
arun-neotrekker:~ arunsriraman$ docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
544fd2b5b674	bridge	bridge	local
790b79d68240	host	host	local
6aaec591a006	none	null	local


Don't want the bridge driver? Remove it by specifying OPTIONS

```
/etc/sysconfig/docker
```

```
OPTIONS="--bridge=none --log-driver=json-file"
```

Recap - Docker network drivers

Driver Features	Bridge / User defined bridge	Host	Overlay	Macvlan / ipvlan
Connectivity	Same host	Same host	Multi-host	Multi-host
Namespace	Separate	Same as host	Separate	Separate
External connectivity	NAT	Use Host gateway	No external connectivity	Uses underlay gateway
Encapsulation	No double encap	No double encap	Double encap using Vxlan	No double encap
Application	North, South external access	Need full networking control, isolation not needed	Container connectivity across hosts	Containers needing direct underlay networking



Part II - Kubernetes Networking

Fundamental requirements


All containers can communicate with all other containers without NAT

All nodes can communicate with all containers (and vice-versa) without NAT

The IP that a container sees itself as is the same IP that others see it as

Kubernetes networking

- Container-to-Container communication
- Pod-to-Pod communication
- Pod-to-Service (cluster internal) communication
- External-to-Service (cluster external) communication




Container CIDR

Service CIDR

Container-to-Container

Pod

Group of one or more containers with shared storage/network

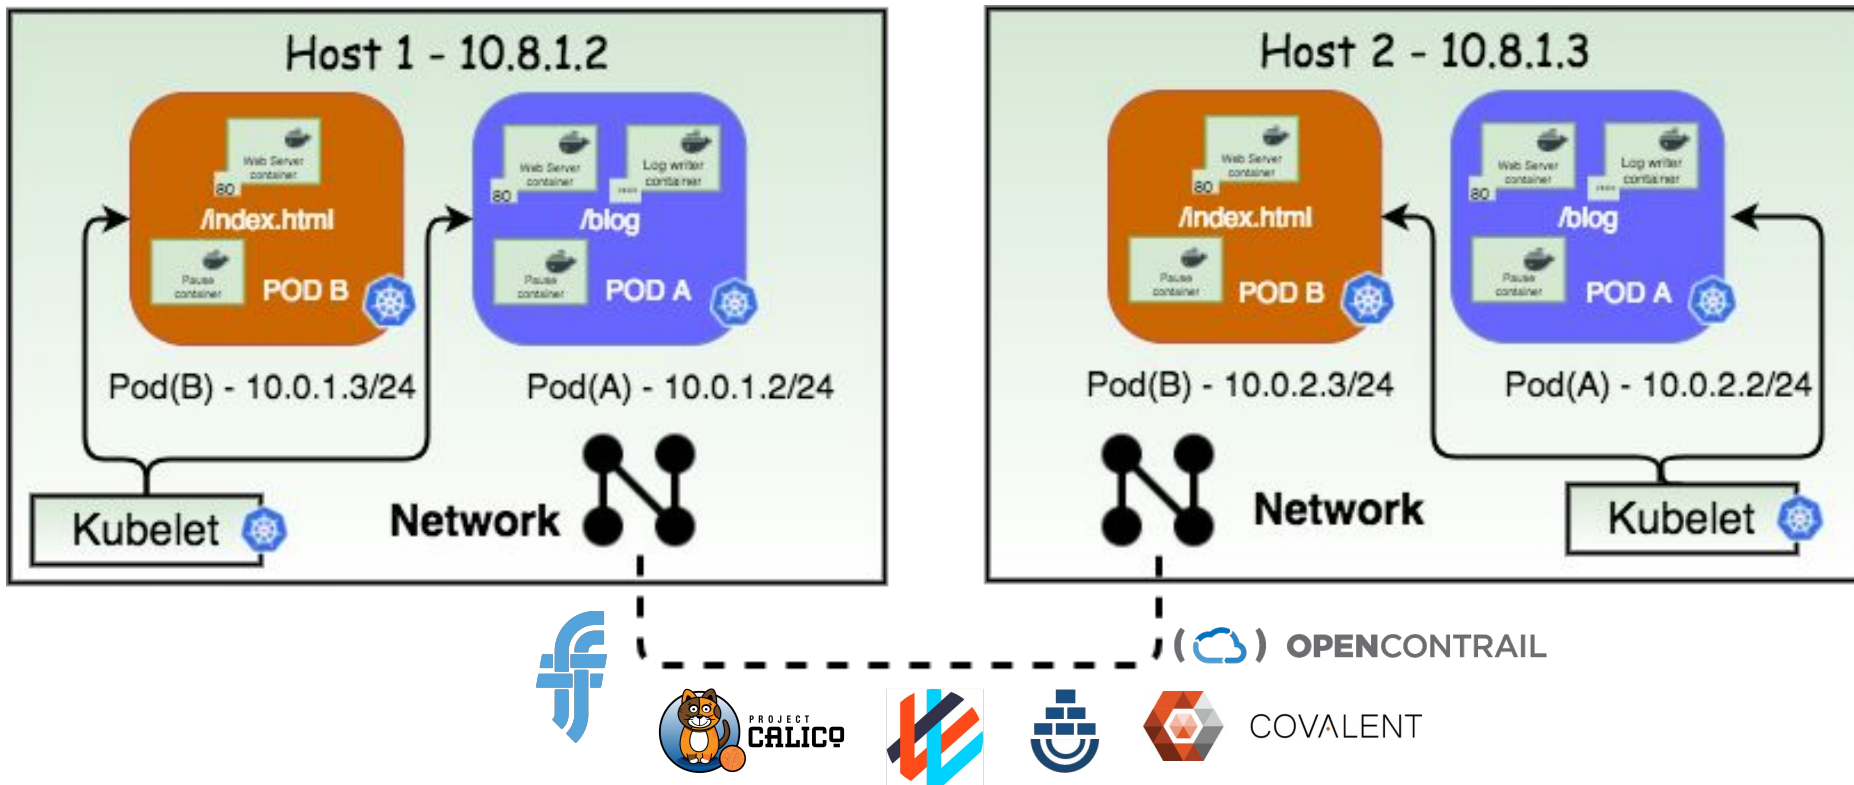


```
[root@ip-10-0-1-25 ~]# docker ps | grep maria
f679a28b57a3      bitnami/mariadb:10.1.23-r2
  "/app-entrypoint.sh /" 2 weeks ago      Up 2 week
s                  k8s_womping-rottweiler-maria
db.7d5c160c_womping-rottweiler-mariadb-155601547-f9lcm_defa
ult_aba49f60-888a-11e7-9059-021eb171b86e_30ebfc40
634acd220b92      gcr.io/google_containers/pause-amd64:3.
0  "/pause" 2 weeks ago      Up 2 week
s                  k8s_POD.d8dbe16c_womping-rot
tweiler-mariadb-155601547-f9lcm_default_aba49f60-888a-11e7-
9059-021eb171b86e_f5e00cb4
```


Container-to-Container takeaways

- Containers in a pod run on the same host.
- A pod generally represents a service unit of an application.
- Uses localhost (127.0.0.1) within the pod's network namespace to communicate with each other
- Containers in the same Pod cannot reuse ports ⚠
- Pause container - Keeps the networking alive
- New concepts: Pod, Pause container

Pod-to-Pod



Pod-to-Pod takeaways

Currently supported networking models -


- Kubenet via kubelet (will be moved out to CNI)
- Multiple network backends via CNI (We'll discuss this in depth later)

Network backend responsible for -

- Pod networking setup
- Pod-to-Pod networking setup (uses L3 BGP like Calico, network overlay like weave, flannel)
- New concepts: Kubelet, CNI, network backend


Kubernetes “Service” Primer

Service - an abstraction which defines a logical set of Pods and a policy by which to access them.



- A service is “generally” backed by pods (endpoints) using a “label selector”.
- Users can explicitly define an endpoint that isn’t backed by pods
- K8s defines many types of services
 - Internal: ClusterIP
 - External: NodePort, LoadBalancer, Ingress

Pod-to-Service (Cluster Internal)




Pod-to-Service takeaways

- Service is a logical definition/collection of pods.
- ClusterIP is allocated from the Services CIDR
- kube-proxy modes
 - userspace
 - iptables (our discussed example)
- New concepts: kube-proxy, kube-dns, Service, clusterIP,

```
Chain KUBE-SVC-GYQQTB6TY565JPRW (1 references)
target      prot opt source                destination
KUBE-SEP-242WNS6JFR3QS6KQ all -- anywhere             anywhere /* default/frontend: */ statistic mode random probability 0.33332999982
KUBE-SEP-3IZ2FS372FZ657HA all -- anywhere             anywhere /* default/frontend: */ statistic mode random probability 0.50000000000
KUBE-SEP-YXDRYNZPYK4TULLG all -- anywhere             anywhere /* default/frontend: */

Chain KUBE-SEP-3IZ2FS372FZ657HA (1 references)
target      prot opt source                destination
KUBE-MARK-MASQ all -- ip-10-49-128-2.us-west-2.compute.internal anywhere /* default/frontend: */
DNAT        tcp  -- anywhere             anywhere /* default/frontend: */ tcp to:10.49.128.2:80
```

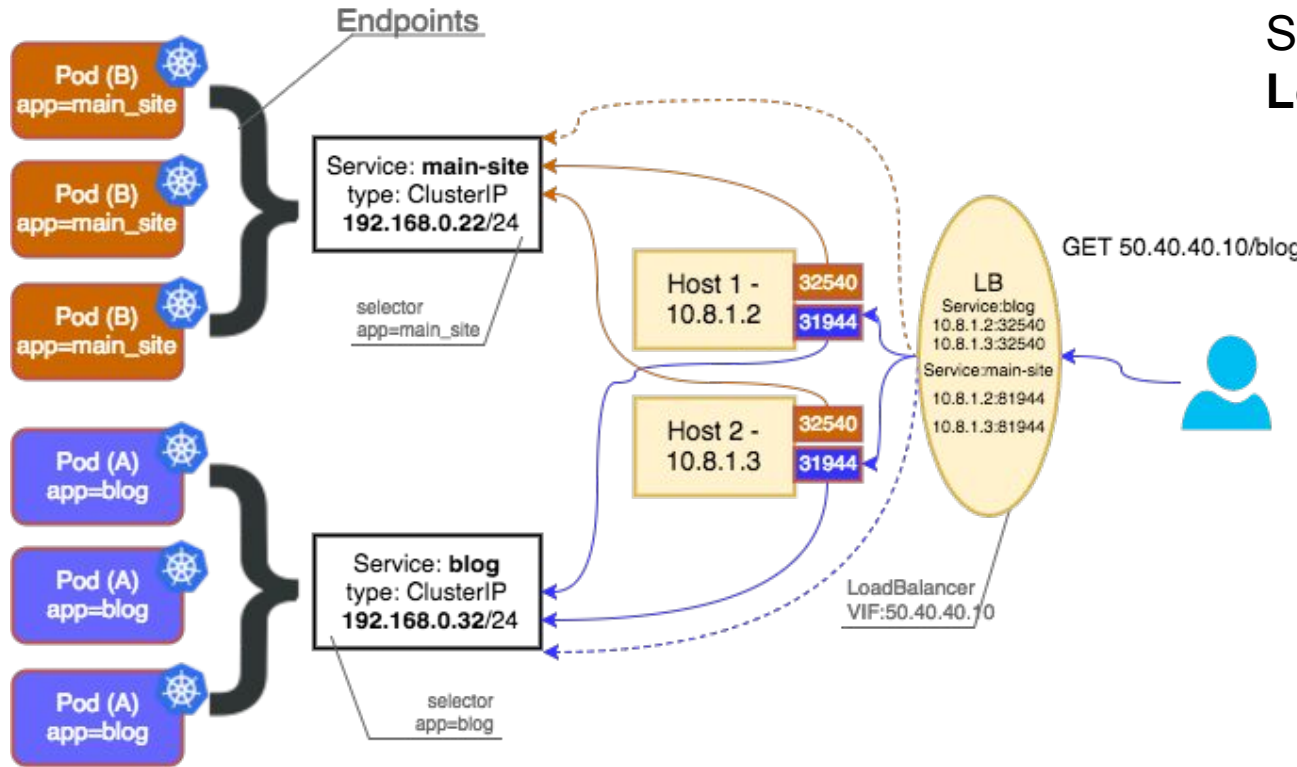
External-to-Service



Service type: **NodePort**

- Kubernetes master allocates a port from a flag-configured range (default: 30000-32767).
- Each Node will proxy that port (the same port number on every Node) into your Service

External-to-Service - II



Service type: **LoadBalancer**

- Fronts the K8s Service
- Traffic from load balancer is directed to backend Pods
- Exactly how that works depends on the cloud provider
- NodePort and ClusterIP to which LB will route are created automatically

External-to-Service III

Ingress

- An Ingress is a collection of rules that allow inbound connections to reach the cluster services.
- Ingress is useful since services typically have internal IPs/endpoints
- All traffic that ends up at an edge router is either dropped or forwarded elsewhere
- Gives services externally-reachable URLs, load balance traffic, terminate SSL, offer name based virtual hosting


External IPs

- A public/external IP points to a node of the cluster
- Service ingresses the requests from the external IP
- Are not managed by K8s

Note: If you came here to understand ingress specifically, let's chat offline. I will cover this if time permits

CNI - Container Network Interface

- Simple interface between container runtime & network
- **CNCF** project. Started by CoreOS for the **rkt** runtime
- Config passed to the NetPlugin by runtime then passed to IPAM
- CNI Interfaces - ADD, DEL



CNI - plugins

CNI Maintained

Plugins that create/delete interfaces

- bridge
- ipvlan
- lo
- macvlan
- vlan
- ptp

IPAM - IP address management

- dhcp - host-local

3rd party/others

- flannel (now under CNI)
- calico
- canal
- weave
- Cilium
- Contrail
- Contiv
- Nuage
- ...

Github repo - <https://github.com/containernetworking/cni>

Thank You

Slides: <https://goo.gl/6XuEnQ>

Code: <https://github.com/sarun87/examples/tree/master/networking>

 [@arun_sriraman](https://twitter.com/arun_sriraman)

Please leave your feedback:
<https://sayat.me/arunsriraman>



Using CNI with individual containers

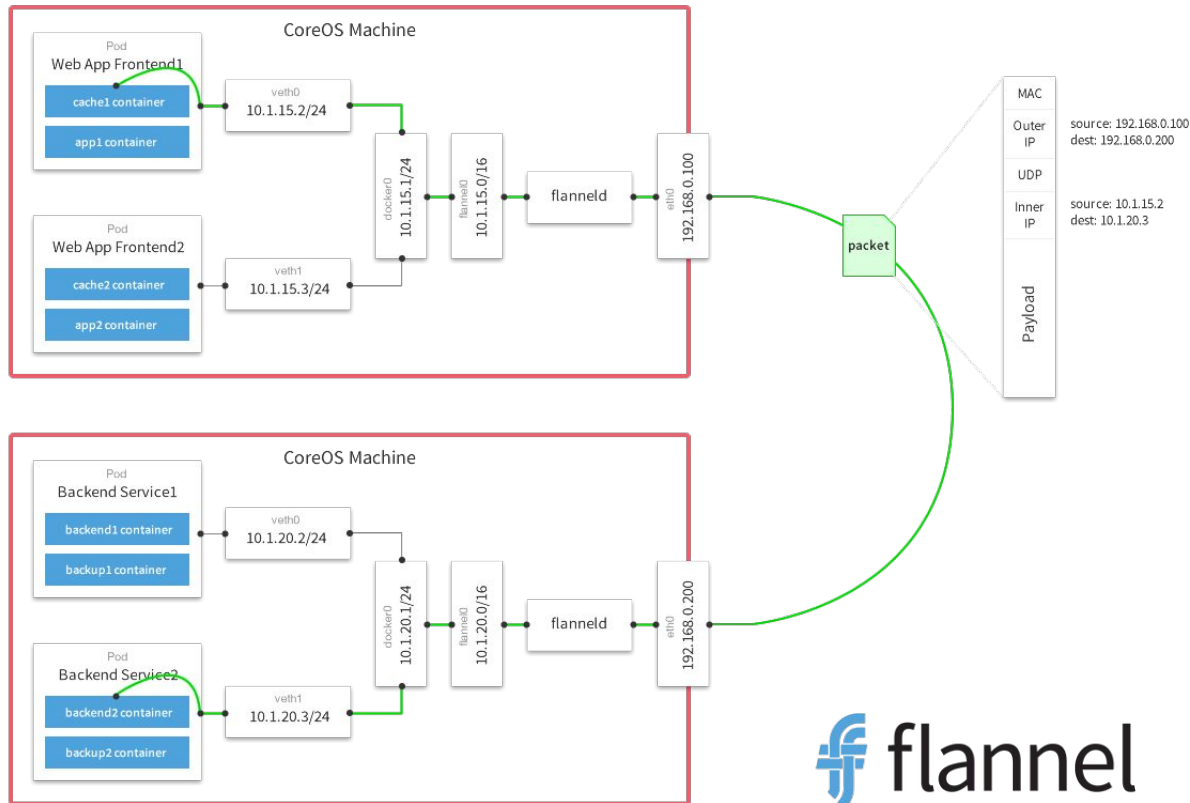
Eg: host-local IPAM. To ADD n/w to a container

```
$ CNI_COMMAND=ADD \  
CNI_CONTAINERID=arun_container_01 \  
CNI_NETNS=/var/run/netns/cni_ipam_eg \  
CNI_IFNAME=eth0 \  
CNI_PATH=/home/ubuntu/cni/bin \  
./host-local < sample_ipam_config
```

```
{  
    "cniVersion": "0.3.1",  
    "ips": [{  
        "version": "4",  
        "address": " 10.10.10.2/24",  
        "gateway": "10.10.10.1"  
    }],  
    "dns": {}  
}
```

```
$ cat sample_ipam_config  
{  
    "cniVersion": "0.3.1",  
    "name": "example-network",  
    "ipam": {  
        "type": "host-local",  
        "subnet": " 10.10.10.0/24",  
        "dataDir":  
        "/home/ubuntu/sample_ipam_datadir"  
    }  
}
```

Flannel network backend

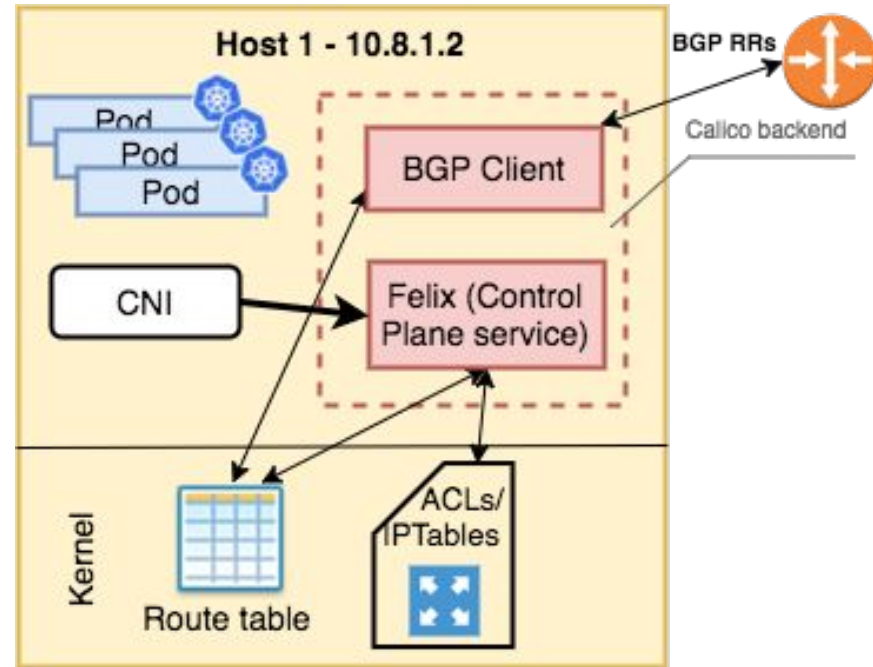


- Uses overlay network for host-host connectivity
- Backends - UDP, vxlan
- flanneld binary runs on every host
- Does **not** perform host - container networking.
- Via CNI, flannel delegates interface operations to bridge driver.

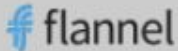




Calico network backend



- Pure L3 based network solution
- Router per node
- Uses BGP
- via CNI plugin - has its own IPAM driver as well
- Supports Kubernetes NetworkPolicy constructs
- BIRD protocol (BGP stack)
- ACL and L3 forwarding performed in the linux kernel
- Ease of debugging
- Scalable



CNI backends summarized

Plugin Features	 flannel	 PROJECT CALICO	 weave	 canal	 Contiv
Main / Networking Plugin	Forwards to bridge driver	Yes	Yes (via bridge plugin)	Yes, bridge driver	Yes
IPAM	host-local	calico-ipam	Weave-ipam / host-local	host-local	<u>Contiv ipam</u>
Host-to-host networking	Overlay - UDP and VXLAN	BGP L3 routing based	Fast data-path and weave router sleeve (VXLAN)	Calico + Flannel	Overlay - VXLAN and VLAN based networks using a vSwitch
K8s <u>NetworkPolicy</u> support	No	Yes	Yes	Yes	Yes
Scalability	Limited	L3 IP. Scalable	Scalable. Fast data-path makes it more efficient	Scalable with advantage of easy setup that flannel brings	Integrates with ACI fabric. Highly scalable with ACI
Debugability	Easy with UDP	Easy since it uses IP	Weave CLI has multiple debugging commands	Mix of calico+flannel	Community and documentation
More to come..					