

# Container and Kubernetes Networking 101

# Before we begin

Slide deck: <https://goo.gl/qmYMPE>

## Poll



**Arun Sriraman**

(Software Engineer)

Platform9 Systems Inc

 @arun\_sriraman  sarun87

# Agenda

## Part I - Container Networking 101

- Why container networking?
- Linux namespaces
- container-to-container communication
- Docker n/w'ing - CNM model
- libnetwork & driver types

## Part II - Kubernetes Networking 101

- K8s networking fundamentals
- Kubernetes communication
  - Container-to-Container
  - Pod-to-Pod
  - Pod-to-Service
  - Service-to-external
- Container Network Interface
- CNI backend (Flannel, Calico)

# **Part I - Container Networking**

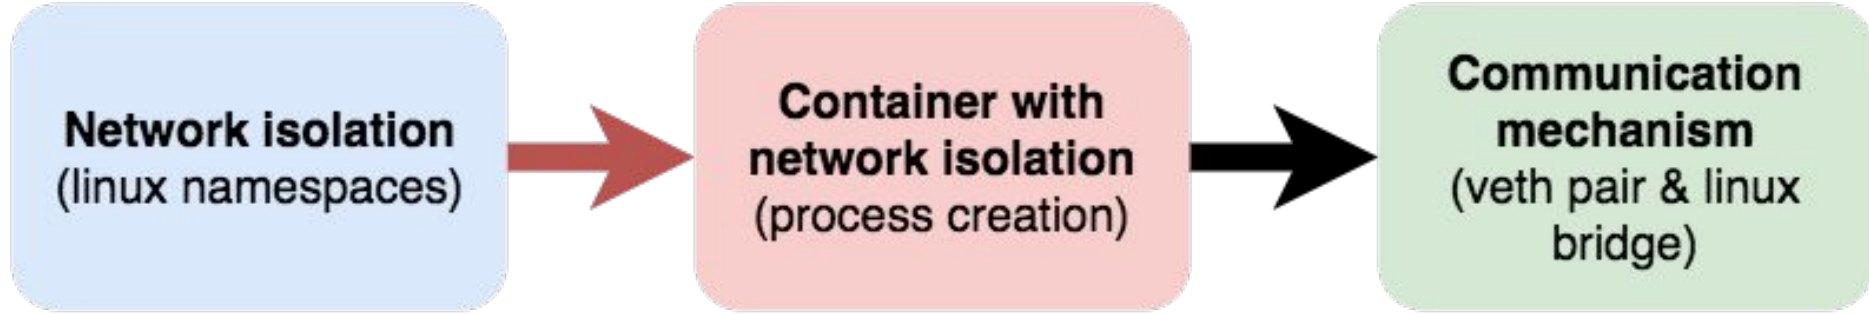
# Why container networking?

- container  $\longleftrightarrow$  container communication
- container  $\longleftrightarrow$  host communication
- user  $\longleftrightarrow$  container communication

Why should we learn the nuts & bolts of the system?

- troubleshooting
- debugging
- building solutions (Service Discovery, Load Balancing)

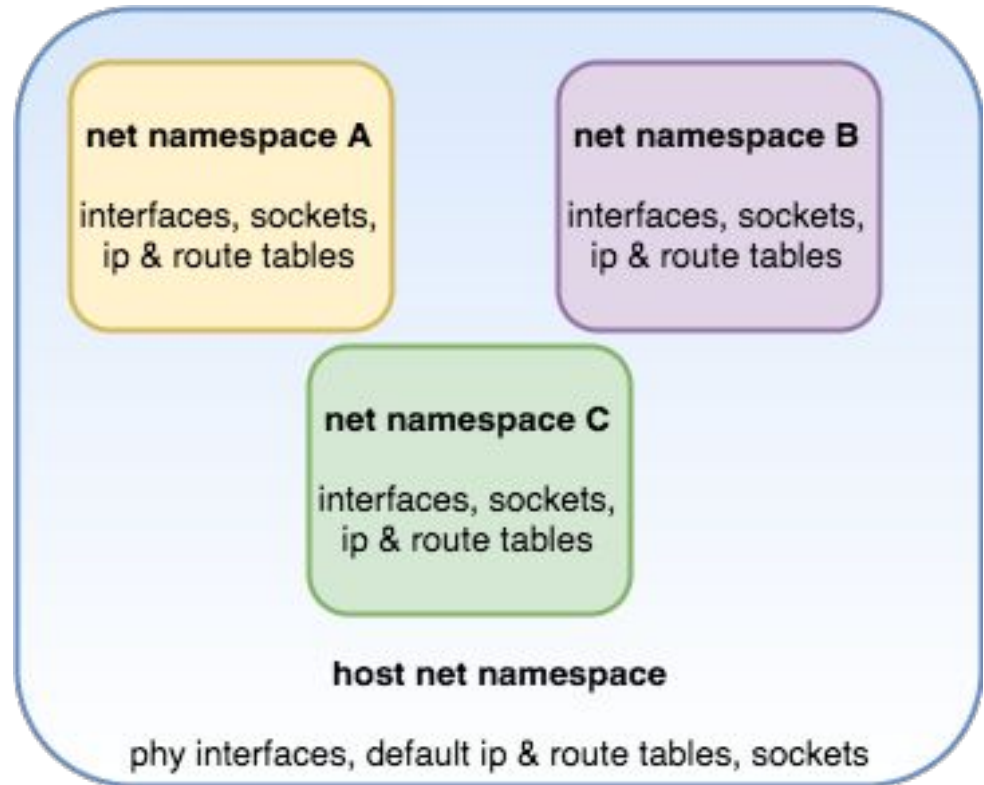
# 3 steps to container networking



# Network Isolation - Network namespaces

A network namespace gets its own private network stack with

- network interfaces (including lo)
- routing tables
- iptables rules
- sockets (ss, netstat)



```
[root@ip-10-0-1-25 ~]# ip netns add A
[root@ip-10-0-1-25 ~]# ip netns add B
[root@ip-10-0-1-25 ~]# ip netns add C
[root@ip-10-0-1-25 ~]# ip netns list
```

Add a network namespace

```
C
B
A
```

List of network namespaces

Namespace A

```
[root@ip-10-0-1-25 ~]# ip netns exec A ip a
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

```
[root@ip-10-0-1-25 ~]#
```

```
[root@ip-10-0-1-25 ~]# ip a
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 9001 qdisc pfifo_fast state UP qlen 1000
    link/ether 02:1e:b1:71:b8:6e brd ff:ff:ff:ff:ff:ff
    inet 10.0.1.25/24 brd 10.0.1.255 scope global dynamic eth0
        valid_lft 2573sec preferred_lft 2573sec
    inet6 fe80::1e:b1ff:fe71:b86e/64 scope link
        valid_lft forever preferred_lft forever
```

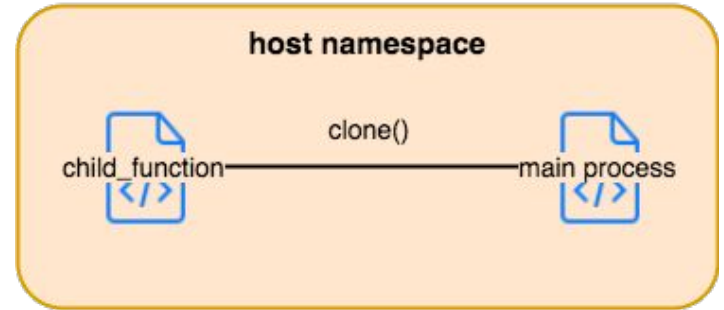
Host network namespace




# Container and isolation

```
flags = CLONE_NEWPID |  
        CLONE_NEWNS | CLONE_NEWNET;  
cpid = clone(child_function,  
            childstack,  
            flags, (void *)argv);
```

without `CLONE_NEWNET` flag

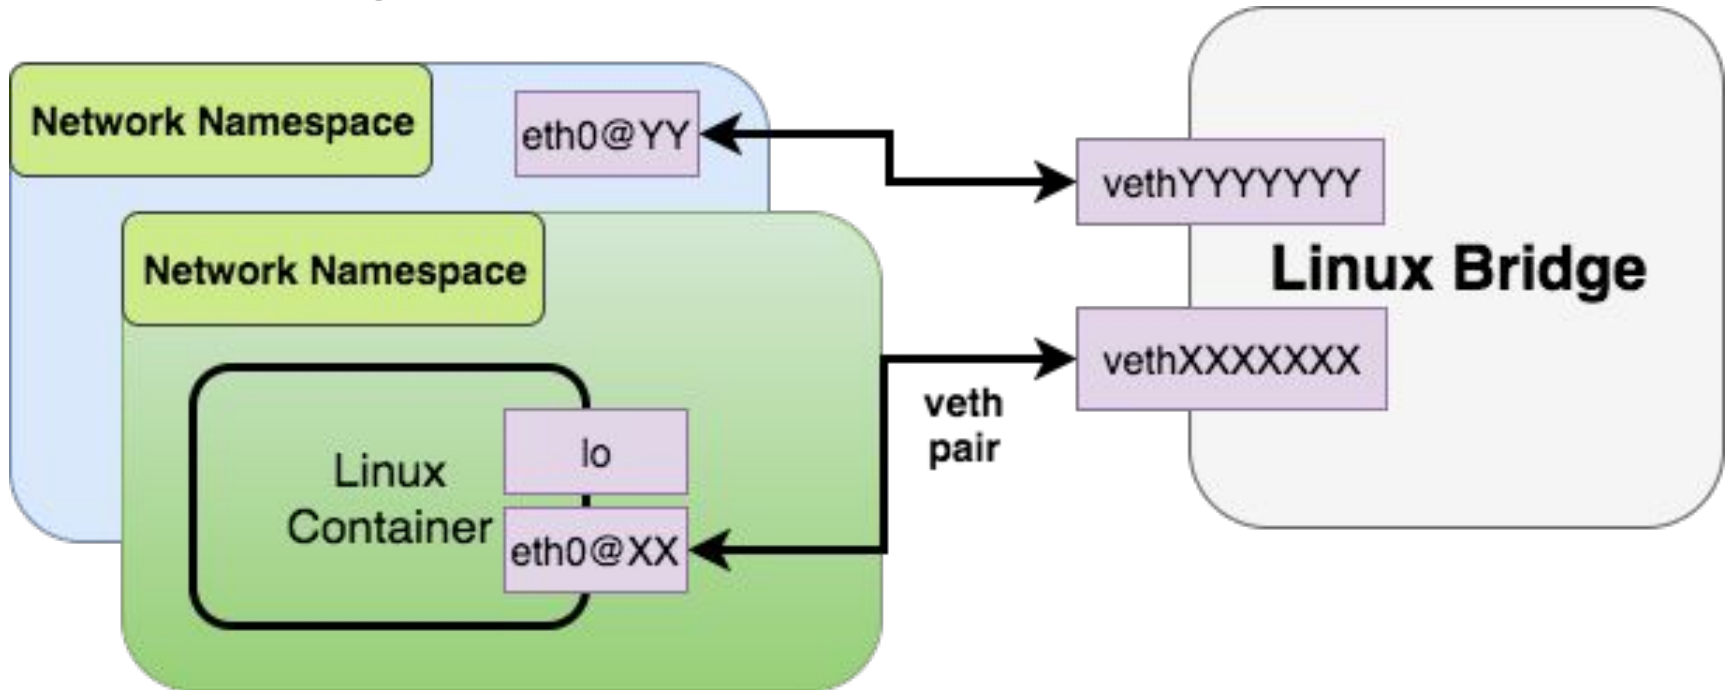


with `CLONE_NEWNET` flag




# Communication mechanism

## Linux bridge and veth pairs

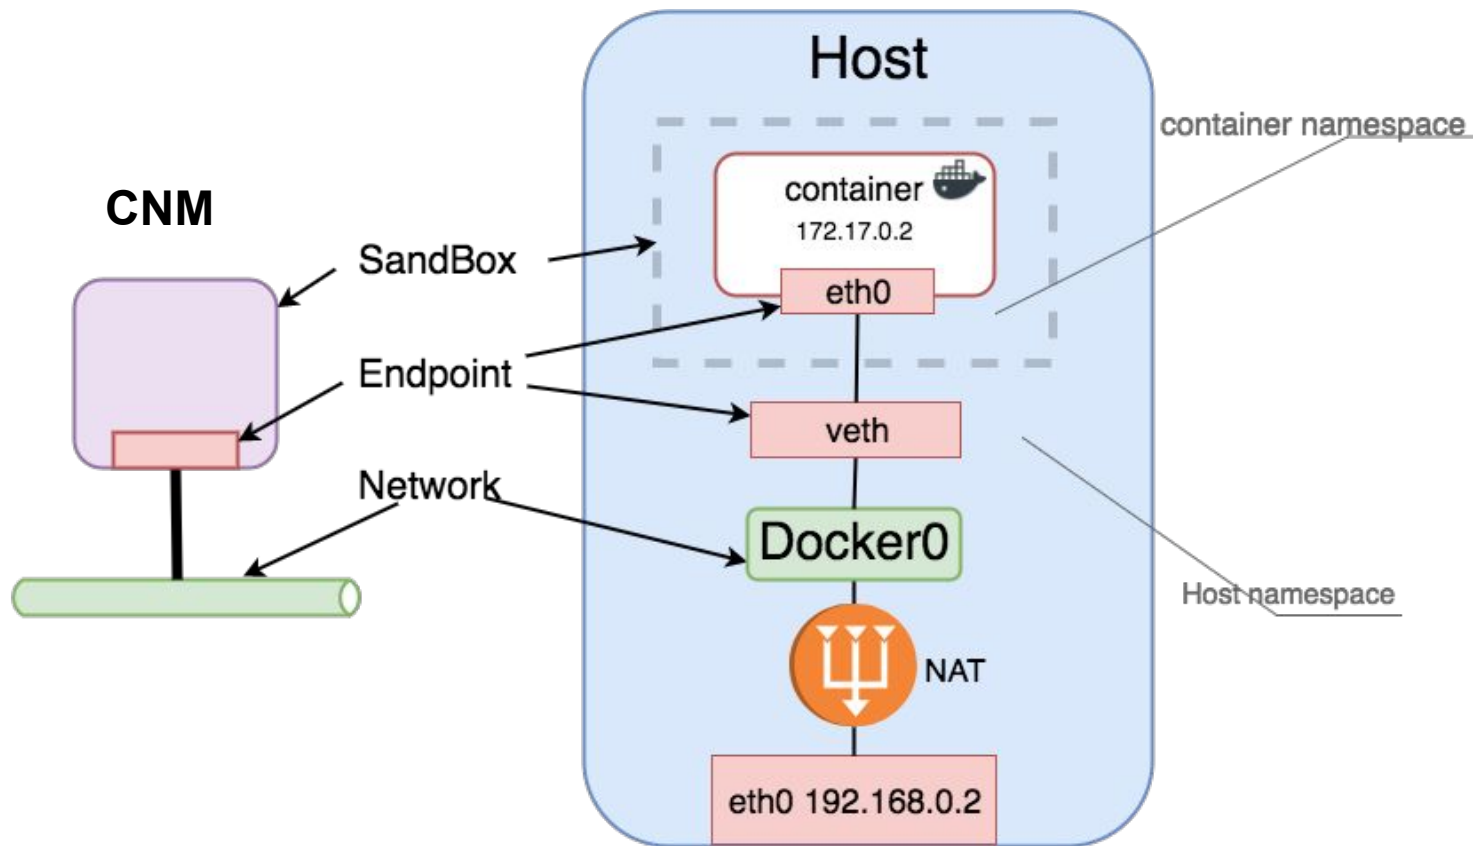


# Container Network Model (CNM)


- Project started by Docker
- Separate networking from container runtime as a library
- Components
  - Sandbox
  - Endpoint
  - Network
- Implemented using libnetwork



# Mapping CNM to Libnetwork (Docker)



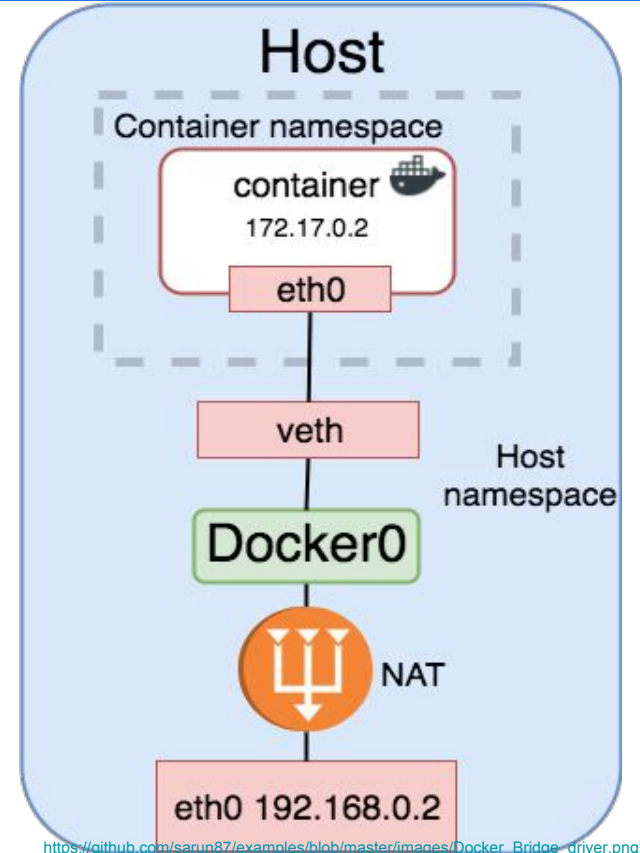
# libnetwork driver types



# Bridge Driver

- Connects docker containers to the network using a veth pair
- Provides out-of-the-box support for bridge based container networking
- Allows creation of user-defined bridges

```
docker network create --driver bridge  
<name>
```

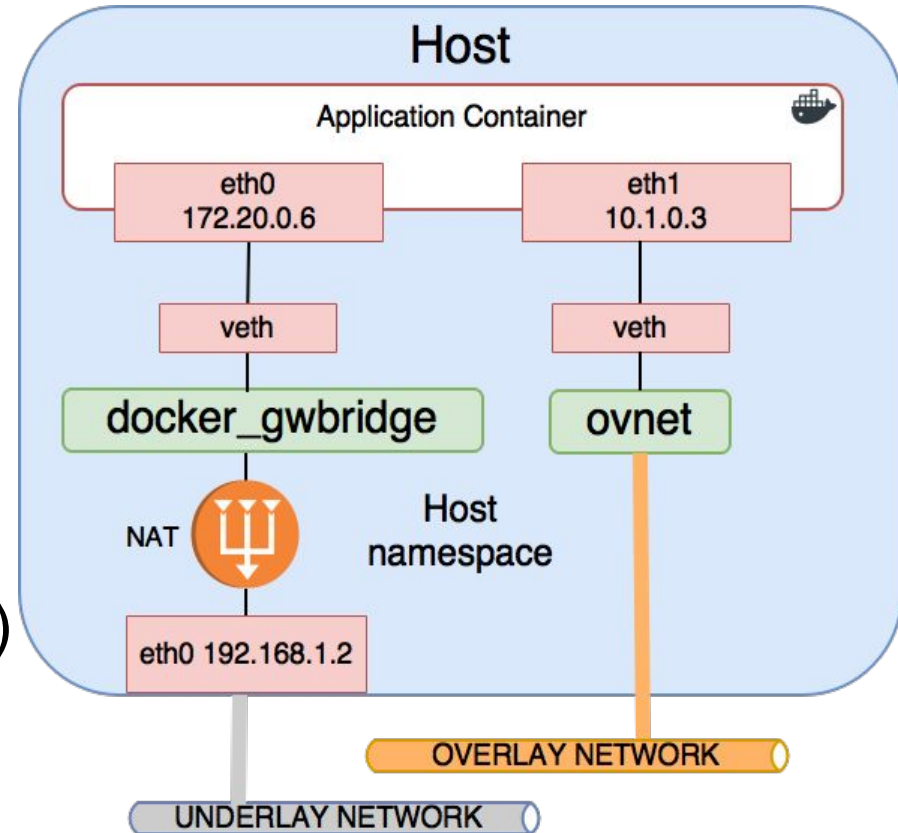


# External Access for Containers

port mapping & NAT



# Overlay Driver

- Multi-host networking
- First-class citizen in docker networking
- Uses swarm-distributed control plane for centralized mgmt, stability & security
- Uses VXLAN encap (decouples container n/w from physical n/w)
- Overlay datapath entirely in kernel space





# Macvlan Driver



# Default Networks Created by Docker

'bridge' using **bridge** driver, 'none' using **null** driver, 'host' using **host** driver

```
arun-neotrekker:~ arunsriraman$ docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
544fd2b5b674	bridge	bridge	local
790b79d68240	host	host	local
6aaec591a006	none	null	local

Don't want the bridge driver? Remove it by specifying OPTIONS

```
/etc/sysconfig/docker
```

```
OPTIONS="--bridge=none --log-driver=json-file"
```

# Recap - Docker network drivers



# **Part II - Kubernetes Networking**

# Fundamental requirements


All containers can communicate with all other containers without NAT

All nodes can communicate with all containers (and vice-versa) without NAT

The IP that a container sees itself as is the same IP that others see it as

# Kubernetes networking

- Container-to-Container communication
- Pod-to-Pod communication
- Pod-to-Service (cluster internal) communication
- External-to-Service (cluster external) communication




Container CIDR

Service CIDR

# Container-to-Container


## Pod

Group of one or more containers with shared storage/network



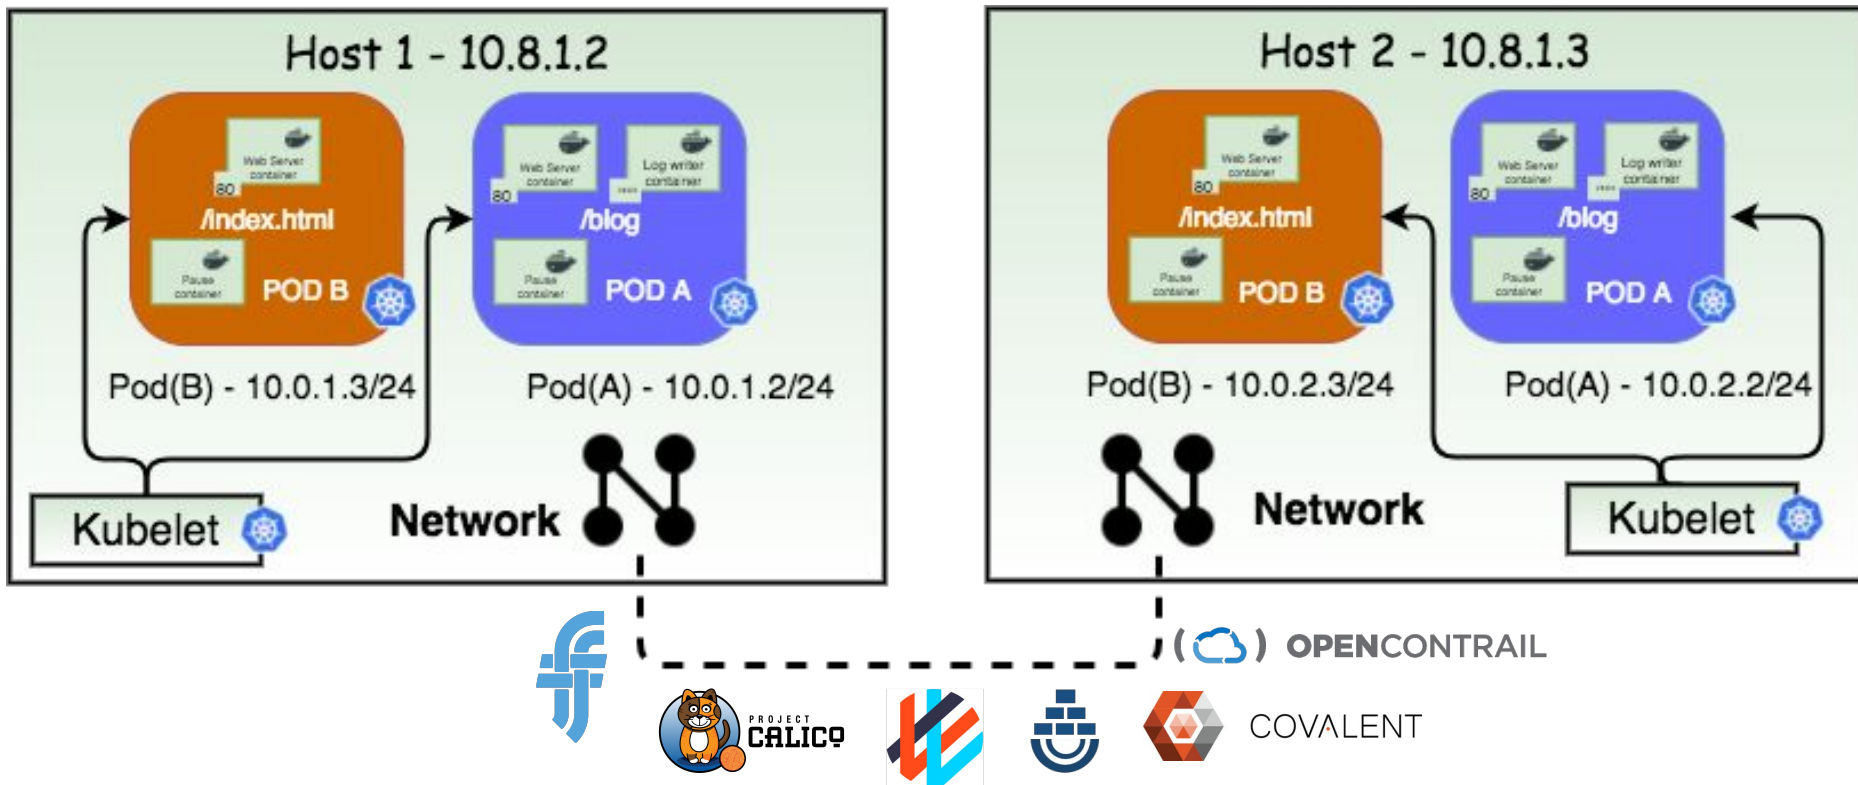
```
[root@ip-10-0-1-25 ~]# docker ps | grep maria
f679a28b57a3      bitnami/mariadb:10.1.23-r2
  "/app-entrypoint.sh /" 2 weeks ago      Up 2 week
s                  k8s_womping-rottweiler-maria
db.7d5c160c_womping-rottweiler-mariadb-155601547-f9lcm_defa
ult_aba49f60-888a-11e7-9059-021eb171b86e_30ebfc40
634acd220b92      gcr.io/google_containers/pause-amd64:3.
0  "/pause" 2 weeks ago      Up 2 week
s                  k8s_POD.d8dbe16c_womping-rot
tweiler-mariadb-155601547-f9lcm_default_aba49f60-888a-11e7-
9059-021eb171b86e_f5e00cb4
```

# Container-to-Container takeaways

- Containers in a pod run on the same host.
- A pod generally represents a service unit of an application.
- Uses localhost (127.0.0.1) within the pod's network namespace to communicate with each other
- Containers in the same Pod cannot reuse ports 
- Pause container - Keeps the networking alive
- New concepts: Pod, Pause container



# Pod-to-Pod



# Pod-to-Pod takeaways

Currently supported networking models -

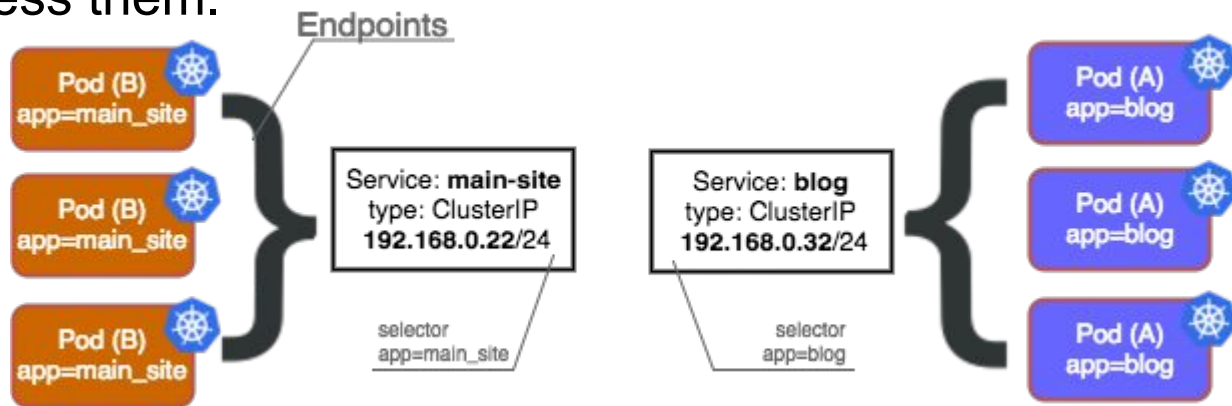
- Kubenet via kubelet (will be moved out to CNI)
- Multiple network backends via CNI (We'll discuss this in depth later)

Network backend responsible for -

- Pod networking setup
- Pod-to-Pod networking setup (uses L3 BGP like Calico, network overlay like weave, flannel)

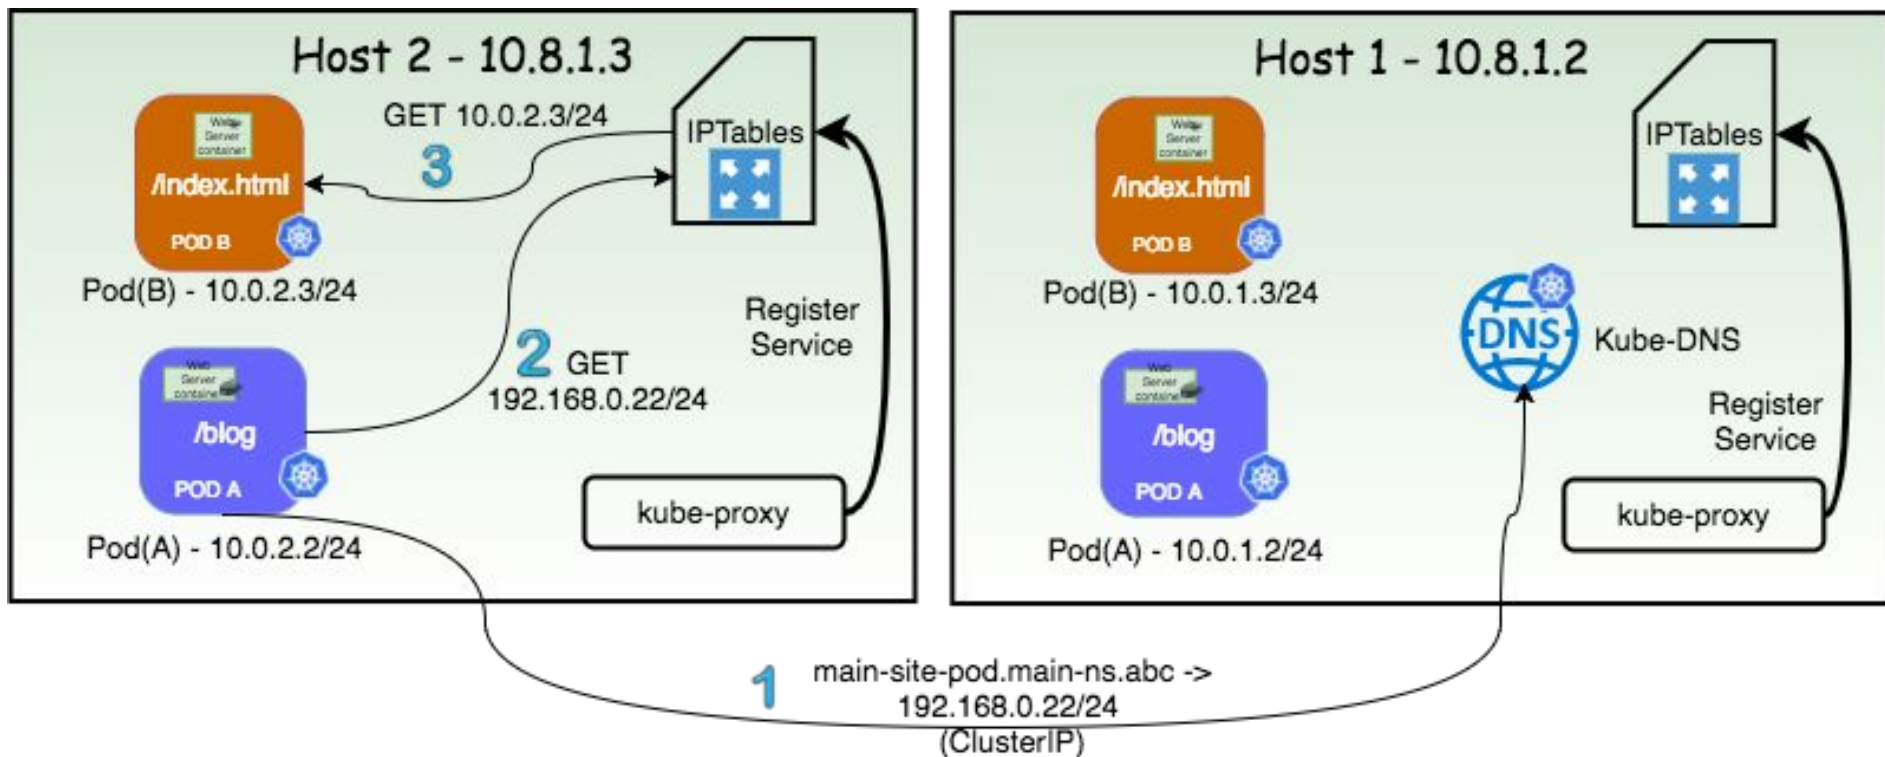
# Kubernetes “Service” Primer

Service - an abstraction which defines a logical set of Pods and a policy by which to access them.



- A service is “generally” backed by pods (endpoints) using a “label selector”.
- Users can explicitly define an endpoint that isn’t backed by pods
- K8s defines many types of services
  - Internal: ClusterIP
  - External: NodePort, LoadBalancer, Ingress

# Pod-to-Service (Cluster Internal)



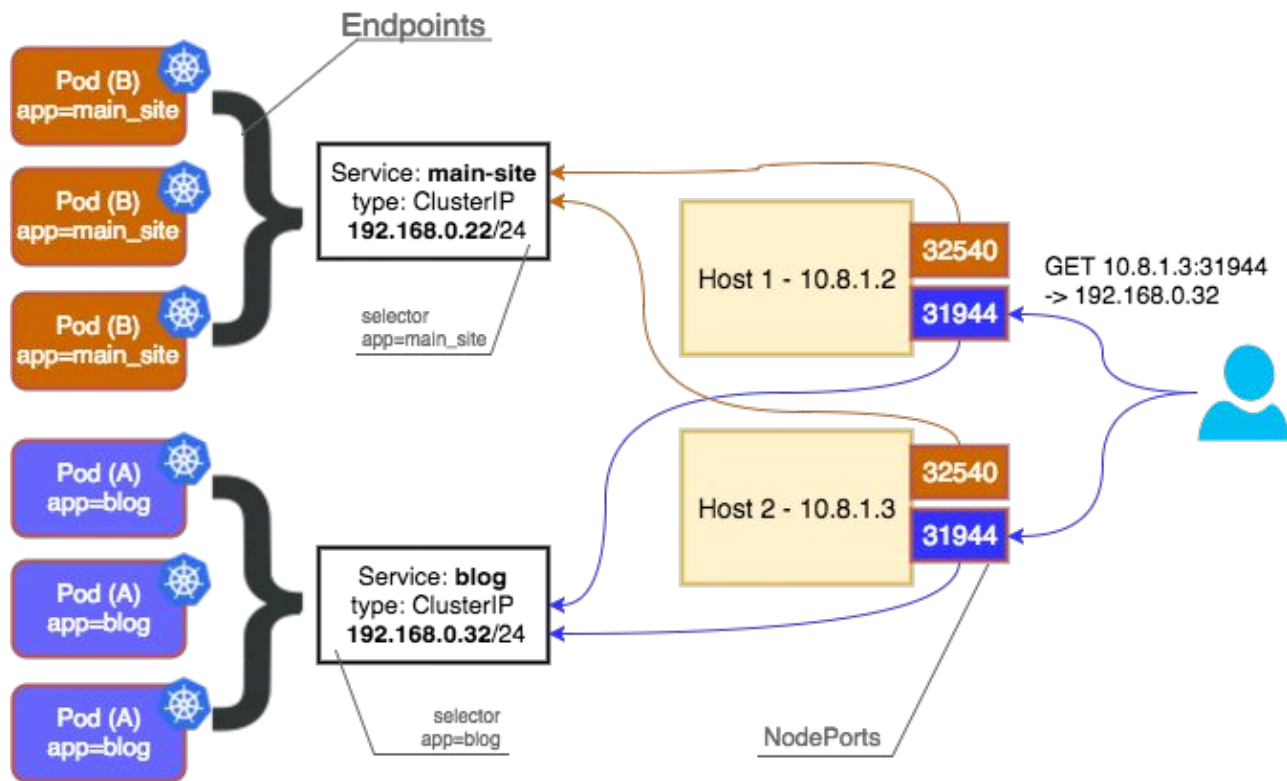
# Pod-to-Service takeaways

- Service is a logical definition/collection of pods.
- ClusterIP is allocated from the Services CIDR
- kube-proxy modes
  - userspace
  - iptables (our discussed example)
- New concepts: kube-proxy, kube-dns, Service, clusterIP,

```
Chain KUBE-SVC-GYQQTB6TY565JPRW (1 references)
target      prot opt source                destination
KUBE-SEP-242WNS6JFR3QS6KQ all -- anywhere             anywhere /* default/frontend: */ statistic mode random probability 0.33332999982
KUBE-SEP-3IZ2FS372FZ657HA all -- anywhere             anywhere /* default/frontend: */ statistic mode random probability 0.50000000000
KUBE-SEP-YXDRYNZPYK4TULLG all -- anywhere             anywhere /* default/frontend: */

Chain KUBE-SEP-3IZ2FS372FZ657HA (1 references)
target      prot opt source                destination
KUBE-MARK-MASQ all -- ip-10-49-128-2.us-west-2.compute.internal anywhere /* default/frontend: */
DNAT        tcp  -- anywhere             anywhere /* default/frontend: */ tcp to:10.49.128.2:80
```


# External-to-Service



Service type: **NodePort**

- Kubernetes master allocates a port from a flag-configured range (default: 30000-32767).
- Each Node will proxy that port (the same port number on every Node) into your Service

# External-to-Service - II



## Service type: **LoadBalancer**

- Fronts the K8s Service
- Traffic from load balancer is directed to backend Pods
- Exactly how that works depends on the cloud provider
- NodePort and ClusterIP to which LB will route are created automatically



# External-to-Service III

## Ingress

- An Ingress is a collection of rules that allow inbound connections to reach the cluster services.
- Ingress is useful since services typically have internal IPs/endpoints
- All traffic that ends up at an edge router is either dropped or forwarded elsewhere
- Gives services externally-reachable URLs, load balance traffic, terminate SSL, offer name based virtual hosting

## External IPs


- A public/external IP points to a node of the cluster
- Service ingresses the requests from the external IP
- Are not managed by K8s

Note: If you came here to understand ingress specifically, let's chat offline. I will cover this if time permits



# CNI - Container Network Interface

- Simple interface between container runtime & network
- **CNCF** project. Started by CoreOS for the **rkt** runtime
- Config passed to the NetPlugin by runtime then passed to IPAM
- CNI Interfaces - ADD, DEL



# CNI - plugins

## CNI Maintained

### Plugins that create/delete interfaces

- bridge
- ipvlan
- lo
- macvlan
- vlan
- ptp

### IPAM - IP address management

- dhcp

## 3rd party/others

- flannel (now under CNI)
- calico
- canal
- weave
- Cilium
- Contrail
- Contiv
- Infoblox
- Romana

Github repo - <https://github.com/containernetworking/cni>

Nuage

# Using CNI with individual containers

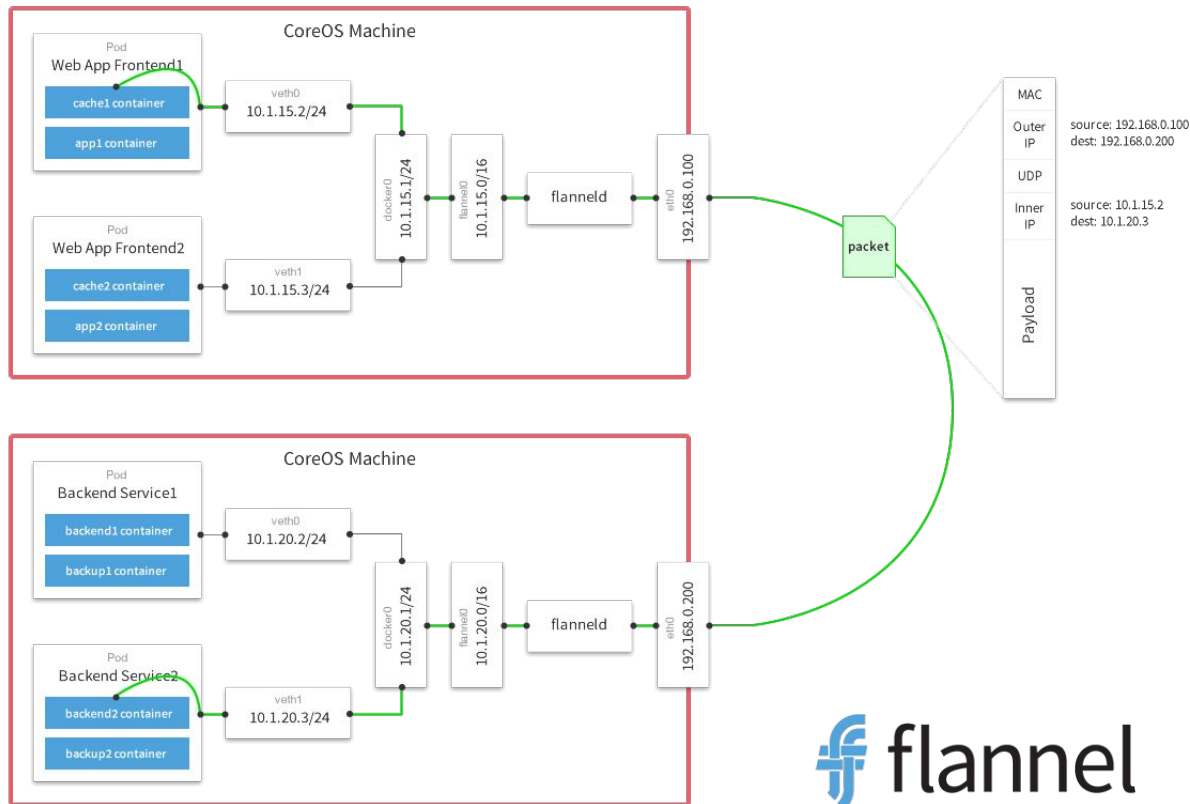
Eg: host-local IPAM. To ADD n/w to a container

```
$ CNI_COMMAND=ADD \  
CNI_CONTAINERID=arun_container_01 \  
CNI_NETNS=/var/run/netns/cni_ipam_eg \  
CNI_IFNAME=eth0 \  
CNI_PATH=/home/ubuntu/cni/bin \  
./host-local < sample_ipam_config
```

```
{  
  "cniVersion": "0.3.1",  
  "ips": [{  
    "version": "4",  
    "address": "10.10.10.2/24",  
    "gateway": "10.10.10.1"  
  }],  
  "dns": {}  
}
```

```
$ cat sample_ipam_config  
{  
  "cniVersion": "0.3.1",  
  "name": "example-network",  
  "ipam": {  
    "type": "host-local",  
    "subnet": "10.10.10.0/24",  
    "dataDir":  
      "/home/ubuntu/sample_ipam_datadir"  
  }  
}
```

# Flannel network backend




- Uses overlay network for host-host connectivity
- Backends - UDP, vxlan
- flanneld binary runs on every host
- Does **not** perform host - container networking.
- Via CNI, flannel delegates interface operations to bridge driver.

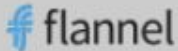




# Calico network backend



- Pure L3 based network solution
- Router per node
- Uses BGP
- via CNI plugin - has its own IPAM driver as well
- Supports Kubernetes NetworkPolicy constructs
- BIRD protocol (BGP stack)
- ACL and L3 forwarding performed in the linux kernel
- Ease of debugging
- Scalable



# CNI backends summarized

Plugin Features	 flannel	 PROJECT CALICO	 weave	 canal	 Contiv
<b>Main / Networking Plugin</b>	Forwards to bridge driver	Yes	Yes (via bridge plugin)	Yes, bridge driver	Yes
<b>IPAM</b>	host-local	calico-ipam	Weave-ipam / host-local	host-local	<u>Contiv ipam</u>
<b>Host-to-host networking</b>	Overlay - UDP and VXLAN	BGP L3 routing based	Fast data-path and weave router sleeve (VXLAN)	Calico + Flannel	Overlay - VXLAN and VLAN based networks using a vSwitch
<b>K8s <u>NetworkPolicy</u> support</b>	No	Yes	Yes	Yes	Yes
<b>Scalability</b>	Limited	L3 IP. Scalable	Scalable. Fast data-path makes it more efficient	Scalable with advantage of easy setup that flannel brings	Integrates with ACI fabric. Highly scalable with ACI
<b>Debugability</b>	Easy with UDP	Easy since it uses IP	Weave CLI has multiple debugging commands	Mix of calico+flannel	Community and documentation
More to come..					

Thank You

**Help me to better help you next time.**

**Questions/Feedback:**

 **@arun\_sriraman**