

# Troubleshooting & Debugging Openstack Neutron

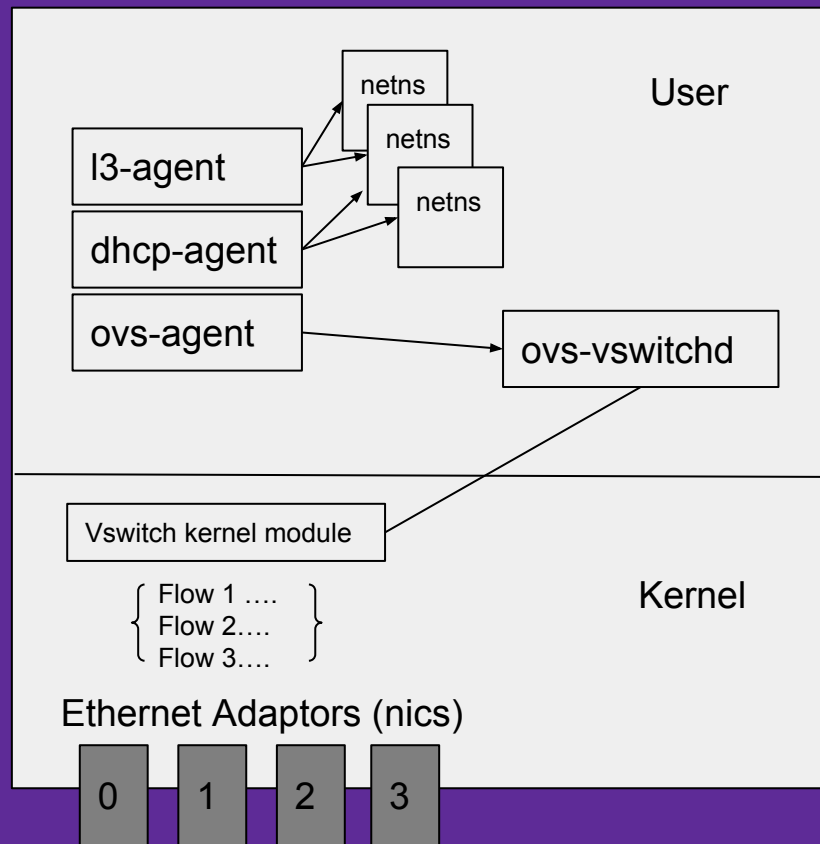
v.1

# Check openstack pre-requisites first

- selinux/apparmor
- Ntp (Network Time Protocol)
- Firewall/networkmanager
- Kernel modules (loaded or not?)
- Sysctl configuration check
- Check if ovs is installed & is running - ovs-vsctl show
- MTU (Max. Transmission Unit)
- Server (physical or virtual?)
- System too slow? Check top/htop, ps aux

# Neutron components

Ovs-agent  
Dhcp-agent  
L3-agent



## Types of Networks supported (Type Drivers)

Flat networks

VLAN

VXLAN tunneling

GRE tunneling

GENEVE tunneling...

...

## Types of L3 architectures

Non DVR (Distributed Virtual Routing)

DVR (Distributed Virtual Routing)

# What to check first... IP connectivity

ip route show

```
[arun@arun-centos71-dev ~]$ ip route show
default via 10.4.0.1 dev ens33 proto static metric 100
10.1.10.3 via 10.4.0.1 dev ens33 proto dhcp metric 100
10.4.0.0/16 dev ens33 proto kernel scope link src 10.4.8.203 metric 100
172.17.0.0/16 dev docker0 proto kernel scope link src 172.17.0.1
192.168.122.0/24 dev virbr0 proto kernel scope link src 192.168.122.1
```

Ping the default gw ip (Eg: 10.4.0.1)

Ping google DNS (8.8.8.8)

Ping google using hostname ([www.google.com](http://www.google.com))

Ping the other host ips. (Both mgmt traffic & data traffic interfaces/ips)

# Check ovs switch connections.

OVS Switches will be connected with each other through patch cables. Check if these cables exist.

```
root@ubuntu01:~# ovs-vsctl show
e8cbdc8e-95f3-46e5-a836-31af58d208b9
```

## Bridge br-ex

```
Port "qg-2dd860c5-f0"
Interface "qg-2dd860c5-f0"
type: internal
```

```
Port br-ex
Interface br-ex
type: internal
```

```
Port phy-br-ex
Interface phy-br-ex
type: patch
options: {peer=int-br-ex}
```

```
Port "eth2"
Interface "eth2"
```

## Bridge br-vlan

```
Port phy-br-vlan
Interface phy-br-vlan
type: patch
options: {peer=int-br-vlan}
```

```
Port "eth1"
Interface "eth1"
Port br-vlan
Interface br-vlan
type: internal
```

## Bridge br-int

```
fail_mode: secure
Port "tap2b68f018-74"
tag: 1
Interface "tap2b68f018-74"
type: internal
Port "qvo69b95a39-b2"
tag: 1
Interface "qvo69b95a39-b2"
```

```
Port int-br-ex
Interface int-br-ex
type: patch
options: {peer=phy-br-ex}
```

```
Port int-br-vlan
Interface int-br-vlan
type: patch
options: {peer=phy-br-vlan}
```

```
Port br-int
Interface br-int
type: internal
Port patch-tun
Interface patch-tun
type: patch
options: {peer=patch-int}
```

## Bridge br-tun

```
fail_mode: secure
Port patch-int
Interface patch-int
type: patch
options: {peer=patch-tun}
Port "vxlan-0a0465fd"
Interface "vxlan-0a0465fd"
type: vxlan
```

```
options: {df_default="true", in_key=flow,
local_ip="10.4.213.120", out_key=flow,
remote_ip="10.4.101.253"}
```

```
Port br-tun
Interface br-tun
type: internal
ovs_version: "2.5.0"
```

# Check if all of the phy ifaces are UP

```
root@ubuntu01:~# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default qlen
1000
    link/ether fa:16:3e:ae:c7:b4 brd ff:ff:ff:ff:ff:ff
    inet 10.4.213.120/16 brd 10.4.255.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::f816:3eff:feae:c7b4/64 scope link
        valid_lft forever preferred_lft forever
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast master ovs-system state UP
group default qlen 1000
    link/ether 00:50:56:8b:9c:3d brd ff:ff:ff:ff:ff:ff
    inet6 fe80::250:56ff:fe8b:9c3d/64 scope link
        valid_lft forever preferred_lft forever
4: eth2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast master ovs-system state UP
group default qlen 1000
    link/ether 00:50:56:8b:af:86 brd ff:ff:ff:ff:ff:ff
    inet6 fe80::250:56ff:fe8b:af86/64 scope link
        valid_lft forever preferred_lft forever
```

# What are the symptoms.....?

1. Instance creation fails
  - a. AMQP Error in nova compute logs.. Oops this could be anything :(
  - b. Port creation fails (neutron logs), check ovs-agent logs. (VIF binding failed? Possibly ovs-agent down.. Nova timeout after 5 mins)
  - c. Retried exception?
2. VM cannot ping (Very helpful information right...?)
  - a. Check if VM has an IP. (Check if VIF binding failed? Probably a DHCP issue? - Try dhclient manually)
  - b. VM has an IP and can't ping another VM on same network? (Try same host)
  - c. VM on same host can't ping each other? (Oops, possibly a security group thing)
  - d. VM on same L2 but different hosts can't ping? (Well well well, is it the underlying network? Try tcpdump - lets take this up later, in the mean time try pinging DHCP server)
  - e. VM on different L2 network (Is router connected, can I ping the router gw interface?)
  - f. VM cannot ping Google. (Can I ping router interface? Can I ping google from router namespace? - lets take a look at this later as well)
  - g. I have a FIP and everything seems fine, just can't get into the VM? (Check security grp)
  - h. I have a FIP and all hell broke loose (VM is kind of dead on the network?) - Advanced debugging...
  - i. ....
3. ....



# Deep dive into network namespaces

Listing namespaces on a linux box:

```
root@ubuntu01:~# ip netns  
snat-203dc92c-18c1-4799-8091-0bcd1f81f3f4  
qrouter-203dc92c-18c1-4799-8091-0bcd1f81f3f4  
Qdhcp-d188b0f8-9f03-4d3a-bf74-9e39de33168e
```

Run commands within a network namespace: `ip netns exec <namespace> <command>`

- NS has its own net stack (ip route, interfaces, iptables ...)

SNAT namespace:

- 2+x ifaces: 1 lo, 1 public network (qg-xx), x tenant n/w lfaces (sg-xx)
- qg-xx interface is part of external bridge and sg-xx is part of integration bridge in the root namespace.
- Ping VM from snat namespace if tenant n/w is connected to the router & ping to 8.8.8.8 does not work.
- Ping External gateway from within this namespace (check external connectivity)

For floating IPs:

- Fip namespace for DVR. For Non-DVR qrouter-xxx namespace.

# Continuing namespaces....

The qdhcp namespace:

- Responsible for per. tenant DHCP.
- Dnsmasq process serves a particular tenant subnet with ip & dns services.
- Dhcp server only responds to IP-mac pair already present in the host file. Does not give out IPs to any dhcp requests.

```
nobody 20506 1 0 Jun06 ? 00:00:00 dnsmasq --no-hosts --no-resolv --strict-order
--except-interface=lo --pid-file=/var/opt/pf9/neutron/dhcp/d188b0f8-9f03-4d3a-bf74-9e39de33168e/pid
--dhcp-hostsfile=/var/opt/pf9/neutron/dhcp/d188b0f8-9f03-4d3a-bf74-9e39de33168e/host
--addn-hosts=/var/opt/pf9/neutron/dhcp/d188b0f8-9f03-4d3a-bf74-9e39de33168e/addn_hosts
--dhcp-optsfile=/var/opt/pf9/neutron/dhcp/d188b0f8-9f03-4d3a-bf74-9e39de33168e/opts
--dhcp-leasefile=/var/opt/pf9/neutron/dhcp/d188b0f8-9f03-4d3a-bf74-9e39de33168e/leases
--dhcp-match=set:ipxe,175 --bind-interfaces --interface=tap2b68f018-74
--dhcp-range=set:tag0,192.168.2.0,static,86400s --dhcp-lease-max=256 --conf-file= --server=8.8.8.8
--domain=openstacklocal
```

All dhcp/dns information you need can be found in the host file. To see if neutron pushed port information (new VMs IP & mac) to the dhcp server, check the hosts file. DHCP uuid is same the network UUID.

```
Eg: root@ubuntu01:~# cat /var/opt/pf9/neutron/dhcp/d188b0f8-9f03-4d3a-bf74-9e39de33168e/host
fa:16:3e:3f:ef:6d,host-192-168-2-3.pf9-demo.com.,192.168.2.3
fa:16:3e:52:10:38,host-192-168-2-17.pf9-demo.com.,192.168.2.17
fa:16:3e:f8:b7:f4,vm03.pf9-demo.com.,192.168.2.21
fa:16:3e:77:24:ab,vm01.pf9-demo.com.,192.168.2.23
```



**KEEP  
CALM  
AND  
TCPDUMP**

# DHCP does not work..... :( what next?

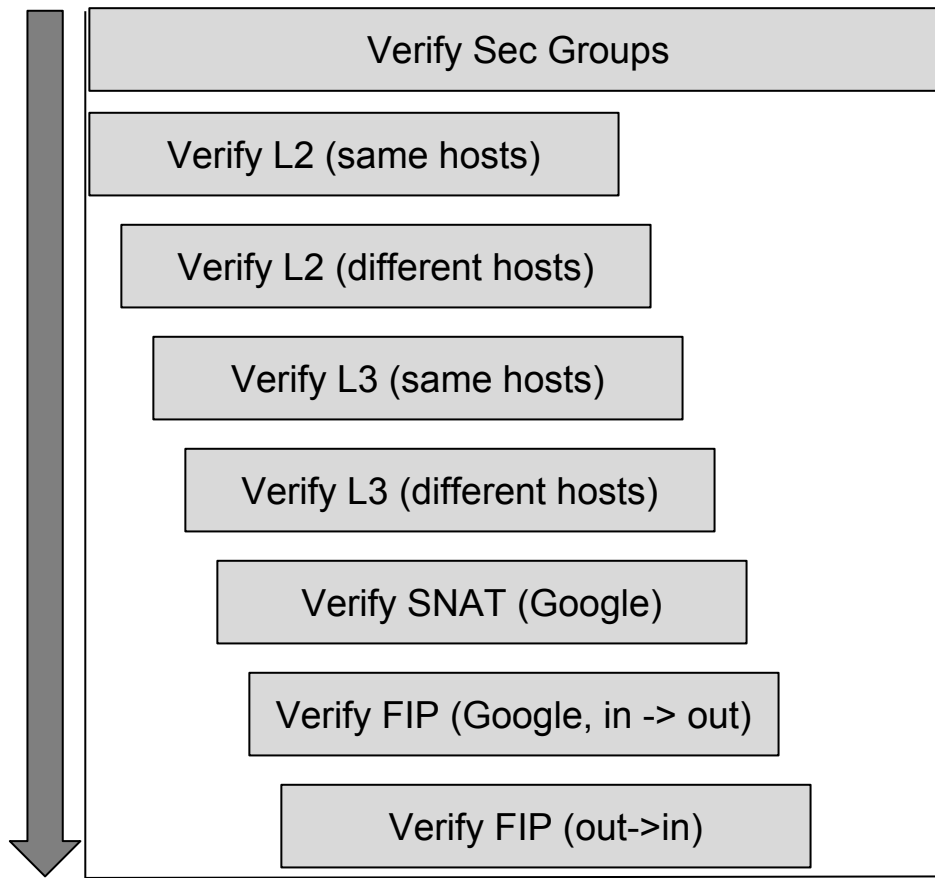
- Try dhclient manually, does it get an IP? (Hope it does..)
- Keep dhclient running. Check if the data interface (phy) can see your bootp request packets
  - `tcpdump -i eth0 -n -e -vv | grep boot`
- Check if the bootp request packet reaches the host running dnsmasq. (tcpdump on phy of host2)
- Check if the packet is seen with
- Repeat: For bootp reply packet

***Find that “blackhole”!***



# Can't figure out the issue? Start from the top!

Arun Sriraman



Great Blog talking about Neutron troubleshooting.

<http://abregman.com/2016/01/06/openstack-neutron-troubleshooting-and-solving-common-problems/>

Avoiding content duplication here :-)

# Retrospect - Most common root causes

- Neutron physical prerequisites setup issues. (Bridge not connected properly. Interface not up)
- Core-network infra configuration (VLANs, upstream bonds, phy interface down)
- MTU problems (hostagent communication, tunneling needs > MTU)
- Arp-cache problem with upstream router disguised as heat stack issue.
- SELinux & firewall problems
- Kernel parameter configuration
- Loop problems in the way of configuring neutron routers. (Education)
- Configuration issues (Need for different ways of networking phy servers)

# Questions

